

Exploring Vision-Based Models for Land-Usage Classification Using Remote Sensing Imagery Data

Presented By - Group 163

Anirudh Kaluri (akaluri)

Rishi Singhal (rsingha4)

Yazhuo Gao (ygao46)

Motivation - Why Land Use Classification?

- Understanding land-use patterns is crucial for sustainable resource management, urban planning, and environmental conservation.
- When combined with deep learning, these images offer insights for disaster recovery, resource allocation, precision agriculture, biodiversity monitoring, and infrastructure planning. A core challenge is accurately classifying land-use patterns from satellite imagery.

Problem Statement

Enhance the understanding of deep learning (DL) models' usefulness for land-use classification using three vision-based neural networks to classify remote sensing images

Dataset Used

- To explore the task of land-use classification, we use the UC-Merced Land-Use Dataset available at Kaggle [1].
- It contains satellite images of different urban regions in the US extracted from USGS National Map Urban Area Imagery collection.

Input & Output of the Task

- Input:

A RGB satellite image representing a region of land use, e.g., a forest. The majority of image size is 256x256.

- Output:

A predicted class label (e.g., forest, river) corresponding to the land-use image feeded to the network.



Prior Work

- Rishi has previously worked with vision based models for different classification tasks.

Methodology


Data Splitting & Preprocessing

- Data splitting - We create a DataLoader which splits train set and validation set with 75% and 25%.
- Data normalization - We resize all images to 256x256 to maintain consistency.
- Transformations - We try multiple combinations of transformations to ensure better training performance.

```
Transform_pipeline = transforms.Compose([
    transforms.Resize((256, 256)), # Resize images
    transforms.RandomHorizontalFlip(p=0.5), # Flip images horizontally with 50% probability (74%)
    transforms.RandomRotation(degrees=30), # Rotate images randomly within ±30 degrees (74%)
    # transforms.RandomResizedCrop(256, scale=(0.8, 1.0)), # Crop images randomly (68%)
    # transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1), # Adjust brightness, contrast, etc. (65%)
    # transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)), # Apply random translation (72%)
    # transforms.RandomPerspective(distortion_scale=0.2, p=0.5), # Apply random perspective transformations (69%)
    transforms.ToTensor(), # Convert image to tensor
    transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5]) # Normalize images
])
```


Baseline CNN Model

- CNN Model - We use a CNN consists of 5 2D-convolution layers with BatchNorm, ReLU, and AvePool. A classifier with dropout is followed at last to handle the features and output (21x1 size). Below is the architecture diagram:



Hyper-parameters & Metrics

- Hyperparameters

Dropout	Learning Rate	Weight Decay	Epochs	Loss function	Optimizer
p = 0.3	3e-4	1e-4	100	Cross Entropy Loss	SGD/Adam

- Evaluation Metrics

Accuracy	Precision	Recall	F1-score
----------	-----------	--------	----------

Since our dataset is fully balanced, accuracy is considered to be the optimal choice to evaluate our model. We also plot **confusion matrix** on the evaluation set. Hence, providing a thorough metrics evaluation in our project.

Results

Analysis Across Multiple Data Transformations

Data Transformation	Accuracy
Random Horizontal Flip	74%
Random Rotation	74%
Random Resized Crop	68%
Color Jitter	65%
Random Affine	72%
Random Perspective	69%

We choose **Random Horizontal Flip and Rotation** as final data transformation techniques.




Adam V/S SGD Optimizer Analysis

Optimizer	Accuracy	Precision	Recall	F1 Score
SGD	83.05%	83.14%	83.65%	82.66%
Adam	80.38%	80.93%	80.99%	80.21%

Using the best 2 data transformations, we analyse which optimizer is best - **SGD performs the best for our task!**

Loss-Accuracy Curves and Confusion Matrix

Using the best hyperparameters we plot the loss-accuracy curves and the confusion matrix (SGD + 2 data transformations + other parameters as discussed previously)



Accuracy - 83.05%, Precision - 83.14%, Recall - 83.65%, F1-score - 82.66%

Note: Since it is overfitting at later epochs we provide metrics for the epoch where we get the best validation accuracy.

Performance Time and Model Size

- Total Trainable Parameters: 920,469
- Model Size: 10.89 MB
- Per Epoch Train Time - 1.6 secs

The model was trained on a NVIDIA A100 GPU.

References

[1] UC Merced Land Use Dataset. Available:
<https://www.kaggle.com/datasets/abdulhasibuddin/uc-merced-land-use-dataset/data>.