

Exploring Vision-Based Models for Land-Usage Classification Using Remote Sensing Imagery Data

Presented By - Group 163

Anirudh Kaluri (akaluri)

Rishi Singhal (rsingha4)

Yazhuo Gao (ygao46)

Motivation - Why Land Use Classification?

- Understanding land-use patterns is crucial for sustainable resource management, urban planning, and environmental conservation.
- When combined with deep learning, these images offer insights for disaster recovery, resource allocation, precision agriculture, biodiversity monitoring, and infrastructure planning. A core challenge is accurately classifying land-use patterns from satellite imagery.

Problem Statement


Enhance the understanding of deep learning (DL) models' usefulness for land-use classification using three vision-based neural networks to classify remote sensing images

Dataset Used

- To explore the task of land-use classification, we use the UC-Merced Land-Use Dataset available at Kaggle [1].
- It contains satellite images of different urban regions in the US extracted from USGS National Map Urban Area Imagery Collection.
- 21,00 RGB satellite images, most of size 256x256 pixels, across 21 land-use classes.

Dataset Used

forest



river




parkinglot



runway



harbor



mobilehome park



denseresidential



overpass




agricultural



freeway




tenniscourt



beach



chaparral



buildings




mediumresidential



sparseresidential




intersection



baseballdiamond




airplane



storagetanks



golfcourse



Input & Output of the Task

- Input:

A RGB satellite image representing a region of land use, e.g., a forest.


- Output:

A predicted class label (e.g., forest, river) corresponding to the land-use image fed into the network.

Image Categories

- 0: agricultural
- 1: airplane
- 2: baseballdiamond
- 3: beach
- 4: buildings
- 5: chaparral
- 6: denseresidential
- 7: forest
- 8: freeway
- 9: golfcourse
- 10: harbor
- 11: intersection
- 12: mediumresidential
- 13: mobilehomepark
- 14: overpass
- 15: parkinglot
- 16: river
- 17: runway
- 18: sparseresidential
- 19: storagetanks
- 20: tenniscourt

forest



river



Prior Work

- Rishi has previously worked with vision based models for different classification tasks.

Methodology


Data Splitting & Preprocessing

- Data splitting - We create a DataLoader which splits train set, validation set and test set - as with 70%, 15% and 15%.
- Data Resize and normalization - We resize all images to 256x256 and normalize to 0 mean & 1 std dev to maintain consistency.
- Transformations - We try multiple combinations of transformations to ensure better training performance.

```
Transform_pipeline = transforms.Compose([
    transforms.Resize((256, 256)), # Resize images
    transforms.RandomHorizontalFlip(p=0.5), # Flip images horizontally with 50% probability (74%)
    transforms.RandomRotation(degrees=30), # Rotate images randomly within ±30 degrees (74%)
    # transforms.RandomResizedCrop(256, scale=(0.8, 1.0)), # Crop images randomly (68%)
    # transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1), # Adjust brightness, contrast, etc. (65%)
    # transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)), # Apply random translation (72%)
    # transforms.RandomPerspective(distortion_scale=0.2, p=0.5), # Apply random perspective transformations (69%)
    transforms.ToTensor(), # Convert image to tensor
    transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5]) # Normalize images
])
```

Baseline CNN Model


- CNN Model - We use a CNN consists of 5 2D-convolution layers with BatchNorm, ReLU, and AvgPool. A classifier with dropout is followed at last to handle the features and output (21x1 size). Below is the architecture diagram:



- Hyperparameters

Dropout Rate	Learning Rate	Weight Decay	Epochs	Loss function	Optimizer
$p = 0.3$	$3e-4$	$1e-4$	100	Cross Entropy Loss	SGD/Adam

ResNet50




We use a 50 layer pre-trained ResNet model.

We finetune this model for our satellite imagery classification task (21 classes).

Learning Rate	Weight Decay	Epochs	Loss function	Optimizer
3e-4	1e-4	100	Cross Entropy Loss	SGD

Vision Transformer - DeiT

Transformer Encoder



We use a 12 layer pre-trained DeiT model (distilled version of ViT-Base).

We finetune this model for our satellite imagery classification task (21 classes).

Learning Rate	Weight Decay	Epochs	Loss function	Optimizer
3e-4	1e-4	50	Cross Entropy Loss	SGD

Evaluation Metrics

- Evaluation Metrics

Accuracy	Precision	Recall	F1-score
----------	-----------	--------	----------

Since our dataset is **fully balanced**, hence **accuracy** is considered to be the optimal choice to evaluate our model. We also plot **confusion matrix** on the test set. Hence, all of these provide a thorough metrics evaluation in our project.

Results - Baseline CNN

Analysis Across Multiple Data Transformations

Data Transformation	Accuracy
Random Horizontal Flip	74%
Random Rotation	74%
Random Resized Crop	68%
Color Jitter	65%
Random Affine	72%
Random Perspective	69%

We choose **Random Horizontal Flip and Rotation** as final data transformation techniques.


Adam V/S SGD Optimizer Analysis

Optimizer	Accuracy	Precision	Recall	F1 Score
SGD	83.05%	83.14%	83.65%	82.66%
Adam	80.38%	80.93%	80.99%	80.21%

Using the best 2 data transformations, we analyse which optimizer is best - **SGD performs the best for our task!**

Loss-Accuracy Curves and Confusion Matrix

Using the best hyperparameters we plot the loss-accuracy curves and the confusion matrix (SGD + 2 data transformations + other parameters as discussed previously)



Test-Set Performance:

Accuracy - 83.05%, Precision - 83.14%, Recall - 83.65%, F1-score - 82.66%

Note: Since it is overfitting at later epochs we use the validation set to find best model, and then use the best model for test-set

Performance Time & Trainable Parameters Count

- Total Trainable Parameters: 920,469
- Per Epoch Train Time - 1.6 secs

All the models are also trained on H100 GPU.

Results - ResNet50

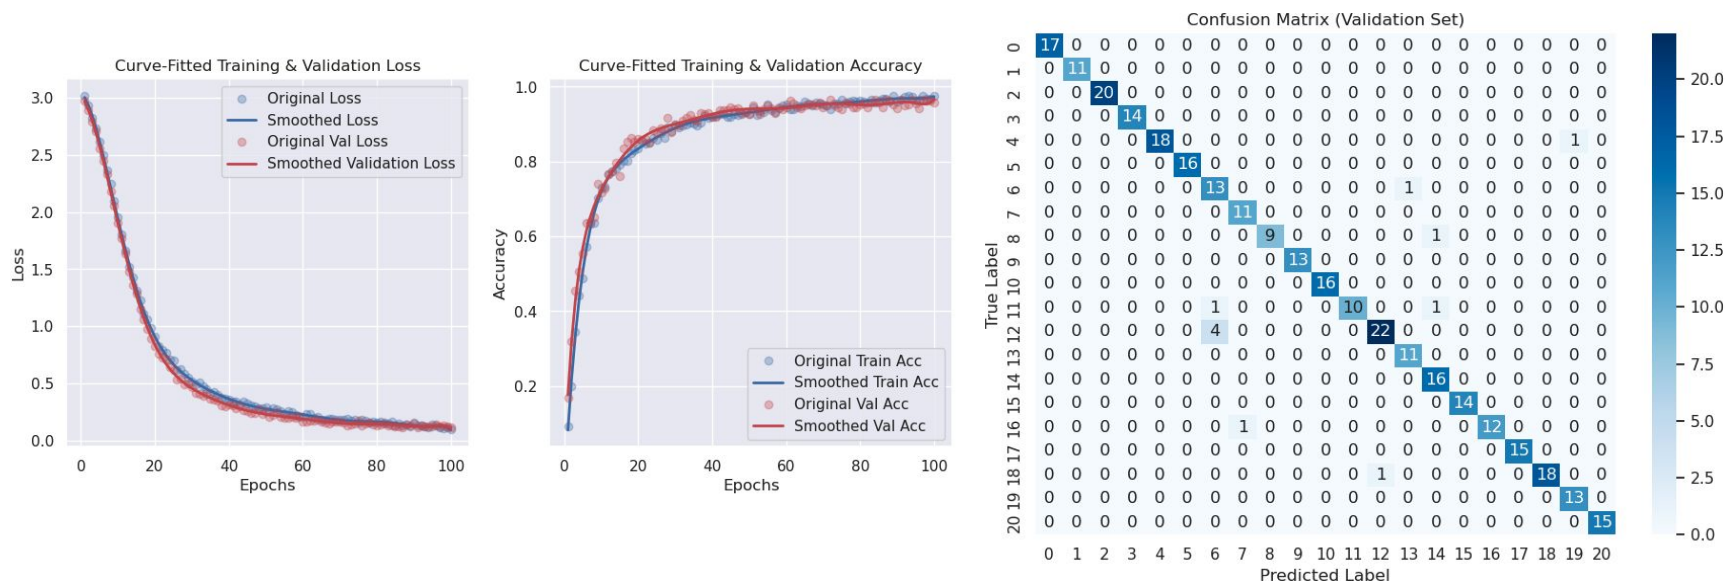
Unfreezing Layers in Group & Trainable Parameters, Train Time per Epoch

Unfrezed (Trainable) Components	Test Accuracy	Total Trainable Parameters	Train Time Per Epoch
All 4 stages + CH	96.51%	23,551,061	2.87s
Last 3 stages + CH	93.02%	23,335,253	2.79s
Last 2 stages + CH	94.29%	22,115,669	2.71s
Last 1 stage + CH	91.43%	15,017,301	7.64s
Only CH	83.17%	52,565	6.83s

- **Unfreezing all layers** give the best performance but also has higher train time and resource utilization. Also, the performance gain is very good (13% up) in comparison to the baseline CNN accuracy.

Loss-Accuracy Curves and Confusion Matrix

Using the best hyperparameters we plot the loss-accuracy curves and the confusion matrix (SGD + 2 data transformations + other parameters as discussed previously) for the completely unfrozen model



Test-Set Performance:

Accuracy - 96.51%, Precision - 96.81%, Recall - 96.79%, F1-score - 96.62%

Note: Since it is overfitting at later epochs we use the validation set to find best model, and then use the best model for test-set

Results - DeiT



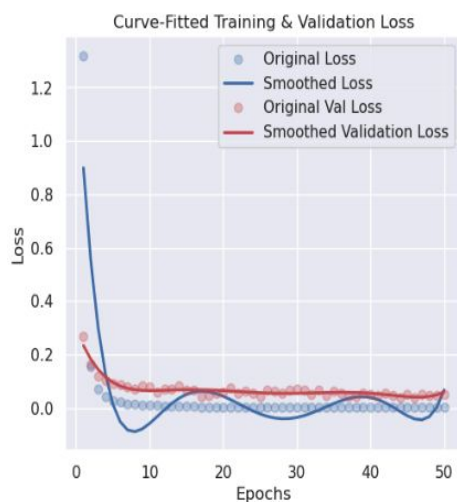
Unfreezing Layers in Group & Trainable Parameters, Train Time per Epoch

Unfrezed (Trainable) Components	Test Accuracy	Total Trainable Parameters	Train Time Per Epoch
All 12 layers + Classification Head (CH)	98.1%	85,816,341	9.7s
Last 9 layers + CH	97.14%	64,552,725	9.02s
Last 6 layers + CH	96.83%	43,289,109	8.36s
Last 3 layers + CH	95.87%	22,025,493	7.64s
Only CH	95.56%	761,877	6.83s

- **Unfreezing all layers** give the best performance but also has higher train time and resource utilization.
- Furthermore, we observe that just by unfreezing classification head (CH) we can achieve almost 96% test-accuracy.

Loss-Accuracy Curves and Confusion Matrix

Using the best hyperparameters we plot the loss-accuracy curves and the confusion matrix (SGD + 2 data transformations + other parameters as discussed previously) for the completely unfrozen model




Test-Set Performance:

Accuracy - 98.1%, Precision - 98.39%, Recall - 98.34%, F1-score - 98.33%

Note: Since it is overfitting at later epochs we use the validation set to find best model, and then use the best model for test-set

Final Analysis

Test Set Performance and Efficiency Comparison across the 3 Models



- ResNet50 and DeiT outperforms the Baseline model significantly (13-15%)
- ResNet50 and DeiT have almost similar performance
- Number of trainable parameters for ResNet50 and train-time per epoch much lesser than that of DeiT.
- **ResNet50 seems to be the ideal model choice.**

References

[1] UC Merced Land Use Dataset. Available:
<https://www.kaggle.com/datasets/abdulhasibuddin/uc-merced-land-use-dataset/data>.